

INSTITUT  
POLYTECHNIQUE  
DE PARIS

# Approximate Pattern Matching

MARIE Nordine & SAVES Marion





INSTITUT  
POLYTECHNIQUE  
DE PARIS

# Hardware



# Hardware

## Root :

CPU :

Intel(R) Xeon(R) Silver 4116

CPU @ 2.10GHz

40 coeurs

GPU :

Quadro P5000 16Go

## Other :

CPU :

Intel(R) Core(TM)

i5-8400 CPU @

2.80GHz



INSTITUT  
POLYTECHNIQUE  
DE PARIS

# Structure du code



# Structure du code

## Entrée - Sortie

En entrée :

```
./apm tailleVariation fichierDNA pattern1obligatoire pattern2 pattern3 ...
```

En sortie :

```
for ( i = 0 ; i < nb_patterns ; i++ )  
{  
    printf( "Number of matches for pattern <%=s>: %d\n",  
           pattern[i], n_matches[i] ) ;  
}
```

# Structure du code

## Main

1 - Boucle pour chaque pattern

2 - Boucle pour chaque lettre de l'ADN

3- Appel levenshtein

```
for ( i = 0 ; i < nb_patterns ; i++ )
{
    int size_pattern = strlen(pattern[i]) ;

    int * column ;

    n_matches[i] = 0 ;

    column = (int *)malloc( (size_pattern+1) * sizeof( int ) ) ;
    if ( column == NULL )
    {
        ...
    }

    for ( j = 0 ; j < n_bytes ; j++ )
    {
        int distance = 0 ;
        int size ;

        if APM_DEBUG ...
        endif

        size = size_pattern ;
        if ( n_bytes - j < size_pattern )
        {
            size = n_bytes - j ;
        }

        distance = levenshtein( pattern[i], &buf[j], size, column ) ;

        if ( distance <= approx_factor ) {
            n_matches[i]++ ;
        }
    }

    free( column );
}
```

# Structure du code

## Levenshtein

Pattern :

AG

ADN:

ACGT

Levenshtein :

ACGT -> AG et AC distance 1

ACGT -> AG et CG distance 1

Ç

ACGT -> AG et GT distance 2

ACGT

Si tolérance 0 -> 0 fois

Si tolérance 1 -> 2 fois


Si tolérance =>2 -> 3 fois

```
int levenshtein(char *s1, char *s2, int len, int * column) {
    unsigned int x, y, lastdiag, olddiag;

    for (y = 1; y <= len; y++)
    {
        column[y] = y;
    }

    for (x = 1; x <= len; x++) {
        column[0] = x;
        lastdiag = x-1 ;
        for (y = 1; y <= len; y++) {
            olddiag = column[y];
            column[y] = MIN3(
                column[y] + 1,
                column[y-1] + 1,
                lastdiag + (s1[y-1] == s2[x-1] ? 0 : 1)
            );
            lastdiag = olddiag;
        }
    }

    return(column[len]);
}
```



INSTITUT  
POLYTECHNIQUE  
DE PARIS

# Conditions de test





# Conditions de test

## DataBase :

chr1\_KI270763v1\_alt.fa

## Distance :

5

## Small Pattern :

32 caractères

## Medium Pattern:

96 caractères (3 fois small)


## Large Pattern :

224 caractères (7 fois small)

## Average:

Moyenne pondérée

= (small + 3medium + 7large)/11



INSTITUT  
POLYTECHNIQUE  
DE PARIS

# OpenMP



# OpenMP

Idées non abouties

## **Idée 1 :**

Chaque thread va s'occuper d'un pattern (stratégie dynamique)

Problème : Déséquilibre se crée

## **Idée 2:**

Chaque thread va s'occuper d'un élément de l'ADN

Problème : Résultat faux, peu d'amélioration → Problème de partage de mémoire

# OpenMP

## Idée retenues

```

/* Timer start */
gettimeofday(&t1, NULL);

for ( i = 0 ; i < nb_patterns ; i++ )
{
    int size_pattern = strlen(pattern[i]) ;

    int * column ;

    n_matches[i] = 0 ;

    column = (int *)malloc( (size_pattern+1) * sizeof( int ) ) ;
    if ( column == NULL )
    {
        fprintf( stderr, "Error: unable to allocate memory for column (%ldB)\n",
            (size_pattern+1) * sizeof( int ) ) ;
        return 1 ;
    }

    for ( j = 0 ; j < n_bytes ; j++ )
    {
        int distance = 0 ;
        int size ;

#ifdef APM_DEBUG
        if ( j % 100 == 0 )
        {
            printf( "Processing byte %d (out of %d)\n", j, n_bytes ) ;
        }
#endif

        size = size_pattern ;
        if ( n_bytes - j < size_pattern )
        {
            size = n_bytes - j ;
        }

        distance = levenshtein( pattern[i], &buf[j], size, column ) ;

        if ( distance <= approx_factor ) {
            n_matches[i]++ ;
        }
    }

    free( column ) ;
}

```

```

195 /* Timer start */
196 gettimeofday(&t1, NULL);
197
198 for ( i = 0 ; i < nb_patterns ; i++ )
199 {
200     int size_pattern = strlen(pattern[i]) ;
201
202     int * column ;
203
204     n_matches[i] = 0 ;
205
206+    tmp_matches = 0 ;
207+    #pragma omp parallel
208+    {
209+        #pragma omp for schedule(guided) reduction(+:tmp_matches)
210+        for ( j = 0 ; j < n_bytes ; j++ )
211+        {
212            column = (int *)malloc( (size_pattern+1) * sizeof( int ) ) ;
213            if ( column == NULL )
214            {
215                fprintf( stderr, "Error: unable to allocate memory for column (%ldB)\n",
216                    (size_pattern+1) * sizeof( int ) ) ;
217+                exit(1);
218            }
219
220            int distance = 0 ;
221            int size ;
222
223+            #if APM_DEBUG
224+                if ( j % 100 == 0 )
225+                {
226+                    printf( "Processing byte %d (out of %d)\n", j, n_bytes ) ;
227+                }
228+            #endif
229
230            size = size_pattern ;
231            if ( n_bytes - j < size_pattern )
232            {
233                size = n_bytes - j ;
234            }
235
236            distance = levenshtein( pattern[i], &buf[j], size, column ) ;
237
238+            if ( distance <= approx_factor ) {
239+                tmp_matches = tmp_matches + 1 ;
240+            }
241+        }
242        free( column ) ;
243+        n_matches[i] = tmp_matches;
244    }

```

# OpenMP

## Idée retenues

```
#include <omp.h>
```

```
int tmp_matches ;
```

```
#pragma omp parallel
{
    #pragma omp for schedule(guided) reduction(+:tmp_matches)
    for ( j = 0 ; j < n_bytes ; j++)
    {
        column = (int *)malloc( (size_pattern+1) * sizeof( int ) ) ;
        if ( column == NULL )
        {
            fprintf( stderr, "Error: unable to allocate memory for column (%ldB)\n",
                    (size_pattern+1) * sizeof( int ) ) ;
            exit(1);
        }
        int distance = 0 ;
        int size ;

        #if APM_DEBUG
        if ( j % 100 == 0 )
        {
            printf( "Procesing byte %d (out of %d)\n", j, n_bytes ) ;
        }
        #endif


        size = size_pattern ;
        if ( n_bytes - j < size_pattern )
        {
            size = n_bytes - j ;
        }

        distance = levenshtein( pattern[i], &buf[j], size, column ) ;

        if ( distance <= approx_factor ) {
            tmp_matches = tmp_matches + 1 ;
        }
    }
}
```


# OpenMP

## Trace




# OpenMP

## SpeedUp




# OpenMP

## SpeedUp







INSTITUT  
POLYTECHNIQUE  
DE PARIS

# MPI



## Fichier d'ADN sur une machine

```
ccaggctggagtgtagtgggtggcactatcacagctcactgcagcctcgat
ctcgtaggctcaagggatcctcctgcctcagcctctcatgtagctgggac
ctgactagctggccagccatatcccgcctaattttatttagttagttagt
tagttagttagttagttagttagttaggaagagacaggtctccccatgtt
ggccaggctggctcctcaactcctgagatcaaatgatcctctcacctcatc
ctcacaagtgtctgggattacaggccagagccacagcaactggccTTGTT
TAGTTTCAGTGAATAAAATTGCCTTTTGAGGGAGCTACAGAGGCAATATA
GCAGCTTGAATCCAGGAAGCTTCCTTGGTCTTGTGTTTCTAAGGGTCCC
AGGCATATTATAAGCAACTGCTCCTGATGGAGGCGCtagtgtatgaagag
ctaggttctggagctagaccgctggacttctgtcttgatcatgtgcttt
tgggaacatgacctatcccatcagtcctcaccctcttcatatgcaaaagt
agaaacaatgacagcagctgcctcgttgggttactatgatgatatatgt
catgacttaaaacacataagggtgccgacagcagtggtgtacgtaatgg
gtacttaacagaGCCATTCTTTTCAGGGTCAATGGGGTGATGGCTATTTT
ATTTTTTCTTTTTTTCAATTTGAGACCAAGGACTCAAGAAAAAAATT
CAggccagacgtagtggctcacgcctctaattcccagcactttggaaggct
gaggtgggtggatctctggaggtcagaagttcaaaaccagcctggccaac
atggagaaacccatctctactaaaaatacaaaaattagccgggcgtggt
ggcacatgcctgtaataccagctactcgggaaactgaggtggcaggatca
cttgaactcaggaggcggaggttgacgtgacctgggatcatgcctctgta
ctcctgcactccagcctgggggaaagaacaagactctgtctaaaaaaaaa
aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaTCATCCTGAAGTCATAAAGTGA
AAGAGTAGCAAAGGTGGGCAGGCAGGAAGTACTCAGCAGTGTAGAATGGG
CCCTAGAACCAAGGACTATTCAGAGTCATCTGAAGTCAAGTCAAGGC
ATCACCAACAGGCCAGTCTGCCCCAGCTTAAGCCCTGATACTAGGGGGG
CCAGGCTGGGTGGGGGCGGGTAGGCGGTGGCCCTTGCAAAACAGCCTCTA
TCCTGTGCAACAGTTTATAATTCAGCAGAGAGTGTTTTGTTTTTAAGAAA
```

## Fichier d'ADN sur trois machines

```
ccaggctggagtgtagtgggtggcactatcacagctcactgcagcctcgat
ctcgtaggctcaagggatcctcctgcctcagcctctcatgtagctgggac
ctgactagctggccagccatatcccgcctaattttatttagttagttagt
tagttagttagttagttagttagttaggaagagacaggtctccccatgtt
ggccaggctggctcctcaactcctgagatcaaatgatcctctcacctcatc
ctcacaagtgtctgggattacaggccagagccacagcaactggccTTGTT
TAGTTTCAGTGAATAAAATTGCCTTTTGAGGGAGCTACAGAGGCAATATA
GCAGCTTGAATCCAGGAAGCTTCCTTGGTCTTGTGTTTCTAAGGGTCCC
AGGCATATTATAAGCAACTGCTCCTGATGGAGGCGCtagtgtatgaagag
```

```
ctaggttctggagctagaccgctggacttctgtcttgatcatgtgcttt
tgggaacatgacctatcccatcagtcctcaccctcttcatatgcaaaagt
agaaacaatgacagcagctgcctcgttgggttactatgatgatatatgt
catgacttaaaacacataagggtgccgacagcagtggtgtacgtaatgg
gtacttaacagaGCCATTCTTTTCAGGGTCAATGGGGTGATGGCTATTTT
ATTTTTTCTTTTTTTCAATTTGAGACCAAGGACTCAAGAAAAAAATT
CAggccagacgtagtggctcacgcctctaattcccagcactttggaaggct
gaggtgggtggatctctggaggtcagaagttcaaaaccagcctggccaac
```

```
atggagaaacccatctctactaaaaatacaaaaattagccgggcgtggt
ggcacatgcctgtaataccagctactcgggaaactgaggtggcaggatca
cttgaactcaggaggcggaggttgacgtgacctgggatcatgcctctgta
ctcctgcactccagcctgggggaaagaacaagactctgtctaaaaaaaaa
aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaTCATCCTGAAGTCATAAAGTGA
AAGAGTAGCAAAGGTGGGCAGGCAGGAAGTACTCAGCAGTGTAGAATGGG
CCCTAGAACCAAGGACTATTCAGAGTCATCTGAAGTCAAGTCAAGGC
ATCACCAACAGGCCAGTCTGCCCCAGCTTAAGCCCTGATACTAGGGGGG
CCAGGCTGGGTGGGGGCGGGTAGGCGGTGGCCCTTGCAAAACAGCCTCTA
TCCTGTGCAACAGTTTATAATTCAGCAGAGAGTGTTTTGTTTTTAAGAAA
```

# MPI

Difficultés : Effets de bord

**Pattern :**  
CTAG

**Sur une machine:**  
AGCTAGCTAGCTAGCTAGCTAGCTAGCTAGCT

**Sur deux machines**  
AGCTAGCTAGCTAGCT AGCTAGCTAGCTAGCT



# MPI

Difficultés : Effets de bord

**Pattern :**  
CTAG

**Sur une machine:**  
AGCTAGCTAGCTAGCTAGCTAGCTAGCTAGCT

**Sur deux machines**  
AGCTAGCTAGCTAGCT AGCTAGCTAGCTAGCT



# MPI

## Initialisation

```
#include <mpi.h>
```

```
int
main( int argc, char ** argv )
{
    /* MPI initialisation
    */
    int nb_nodes;
    int rank;
    MPI_Status status;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &nb_nodes);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    #if APM_DEBUG
        char hostname[256];
        gethostname(hostname, sizeof(hostname));
        printf("Process MPI rank %d of PID %d on %s ready for attach\n",rank, getpid(), hostname);
    #endif

    char ** pattern ;
    char * filename ;
    int approx_factor = 0 ;
    int nb_patterns = 0 ;
    int i, j ;
    char * buf ;
    struct timeval t1, t2;
    double duration ;
    int n_bytes ;
    int * n_matches ;
    int * global_matches;
```



# MPI

## Main

```

/* Grab the patterns */
for ( i = 0 ; i < nb_patterns ; i++ )
{
    int l ;

    l = strlen(argv[i+3]) ;
    if ( l <= 0 )
    {
        fprintf( stderr, "Error while parsing argument %d\n", i+3 ) ;
        return 1 ;
    }

    pattern[i] = (char *)malloc( (l+1) * sizeof( char ) ) ;
    if ( pattern[i] == NULL )
    {
        fprintf( stderr, "Unable to allocate string of size %d\n", l ) ;
        return 1 ;
    }

    strncpy( pattern[i], argv[i+3], (l+1) ) ;
}

printf( "Approximate Pattern Mathing: "
        "looking for %d pattern(s) in file %s w/ distance of %d\n",
        nb_patterns, filename, approx_factor ) ;

buf = read_input_file( filename, &n_bytes ) ;
if ( buf == NULL )
{
    return 1 ;
}

/* Allocate the array of matches */
n_matches = (int *)malloc( nb_patterns * sizeof( int ) ) ;
if ( n_matches == NULL )
{
    fprintf( stderr, "Error: unable to allocate memory for %ldB\n",
            nb_patterns * sizeof( int ) ) ;
    return 1 ;
}

163 /* Grab the patterns */
164+ int max_len_pattern = 0;
165 for ( i = 0 ; i < nb_patterns ; i++ )
166 {
167     int l ;

168     l = strlen(argv[i+3]) ;
169     if ( l <= 0 )
170     {
171         fprintf( stderr, "Error while parsing argument %d\n", i+3 ) ;
172         return 1 ;
173     } else if ( l > max_len_pattern )
174     {
175+         max_len_pattern = l;
176     }
177 }
178+
179 pattern[i] = (char *)malloc( (l+1) * sizeof( char ) ) ;
180 if ( pattern[i] == NULL )
181 {
182     fprintf( stderr, "Unable to allocate string of size %d\n", l ) ;
183     return 1 ;
184 }
185
186 strncpy( pattern[i], argv[i+3], (l+1) ) ;
187 }
188
189
190
191 printf( "Approximate Pattern Mathing: "
192         "looking for %d pattern(s) in file %s w/ distance of %d\n",
193         nb_patterns, filename, approx_factor ) ;
194
195+ /* Allocate the array of local matches */
196+ n_matches = (int *)malloc( nb_patterns * sizeof( int ) ) ;
197+ if ( n_matches == NULL )
198+ {
199+     fprintf( stderr, "Error: unable to allocate memory for %ldB\n",
200+             nb_patterns * sizeof( int ) ) ;
201+     return 1 ;
202+ }
203+
204+ /* Allocate the array of global matches for the reductin of local n_matches*/
205+ global_matches = (int *)malloc( nb_patterns * sizeof( int ) ) ;
206+ if ( global_matches == NULL )
207+ {
208+     fprintf( stderr, "Error: unable to allocate memory for %ldB\n",
209+             nb_patterns * sizeof( int ) ) ;
210+     return 1 ;
211+ }
212+

```

# MPI

## Main

```
/* Timer start */
gettimeofday(&t1, NULL);

/* Since we have nb_nodes MPI process we are going to divide our textfile into
nb_nodes parts while taking care that the biggest pattern have access to all
it needs.

rank 0 treats from 0 to n_bytes//nb_nodes - 1 + (max_len_pattern - 1)
rank 1 treats from n_bytes//nb_nodes to 2*(n_bytes//size) - 1 + (max_len_pattern - 1)
...
rank i treat from i*(n_bytes//nb_nodes) to (i+1)*(n_bytes//size) - 1 + (max_len_pattern - 1)
...
rank (nb_nodes-1) treat from (nb_nodes-1)*(n_bytes//nb_nodes) to END
*/

/* rank 0 play the role of divider */
int part_bytes; // the number of bytes of the process part textfile
MPI_Request requests[nb_nodes-1];
MPI_Status statutes[nb_nodes-1];
if (rank == 0) {
    /* reading input file */
    buf = read_input_file( filename, &n_bytes );
    if ( buf == NULL )
    {
        return 1 ;
    }

    int start = 0; // start index of process
    int end = n_bytes/nb_nodes - 1 + (max_len_pattern - 1); // end index of process
    #if APM_DEBUG
        printf( "MPI rank 0 will treat from bytes %d to %d\n", start, end);
    #endif
}
```

```
for (int i = 1; i < nb_nodes; i++) {
    /* Index and process part bytes */
    start += (n_bytes/nb_nodes);
    end += (n_bytes/nb_nodes);
    if (i == nb_nodes - 1 || end > n_bytes) {
        end = n_bytes;
    }
    part_bytes = end - start + 1 ;
    #if APM_DEBUG
        printf("MPI rank %d will treat from bytes %d to %d\n",i,start,end);
    #endif
    /* Sending to each process other than 0*/
    /* the part_bytes so they know how much memory to allocate */
    MPI_Send(&part_bytes,1,MPI_INTEGER,i,0,MPI_COMM_WORLD);
    #if APM_DEBUG
        printf("Rank 0 sended part_bytes : %d to rank %d\n",part_bytes,i);
    #endif
    MPI_Send(&buf[start],part_bytes,MPI_BYTE,i,1,MPI_COMM_WORLD);
    #if APM_DEBUG
        printf("Rank 0 sended a part_buffer to rank %d\n",i);
    #endif
    /* the start & end index of their part */
}
```

```
// Reset part_bytes for process 0 :
part_bytes = n_bytes/nb_nodes - 1 + max_len_pattern - 1;
} else {
    // other process receive :
    // first : part_bytes :
    MPI_Recv(&part_bytes,1,MPI_INTEGER,0,0,MPI_COMM_WORLD,&status);
    // so they know how much to allocate
    buf = (char *) malloc((part_bytes+1)*sizeof(char));
    if ( buf == NULL )
    {
        fprintf( stderr, "Unable to allocate %ld byte(s) for buf array\n",part_bytes);
        return -1;
    }
    // secondly : part textfile :
    MPI_Recv(buf,part_bytes,MPI_BYTE,0,1,MPI_COMM_WORLD,&status);
}
```

# MPI

## Boucle principale

```
for ( i = 0 ; i < nb_patterns ; i++ )
{
    int size_pattern = strlen(pattern[i]) ;

    int * column ;

    n_matches[i] = 0 ;

    column = (int *)malloc( (size_pattern+1) * sizeof( int ) ) ;
    if ( column == NULL )
    {
        fprintf( stderr, "Error: unable to allocate memory for column (%ldB)\n",
                (size_pattern+1) * sizeof( int ) ) ;
        return 1 ;
    }

    for ( j = 0 ; j < n_bytes ; j++ )
    {
        int distance = 0 ;
        int size ;

#ifdef APM_DEBUG
        if ( j % 100 == 0 )
        {
            printf( "Processing byte %d (out of %d)\n", j, n_bytes ) ;
        }
#endif

        size = size_pattern ;
        if ( n_bytes - j < size_pattern )
        {
            size = n_bytes - j ;
        }

        distance = levenshtein( pattern[i], &buf[j], size, column ) ;

        if ( distance <= approx_factor ) {
            n_matches[i]++ ;
        }
    }

    free( column ) ;
}

for ( i = 0 ; i < nb_patterns ; i++ )
{
    int size_pattern = strlen(pattern[i]) ;

    int * column ;

    n_matches[i] = 0 ;

    column = (int *)malloc( (size_pattern+1) * sizeof( int ) ) ;
    if ( column == NULL )
    {
        fprintf( stderr, "Error: unable to allocate memory for column (%ldB)\n",
                (size_pattern+1) * sizeof( int ) ) ;
        return 1 ;
    }

    int j_end = (rank == nb_nodes-1) ? part_bytes : part_bytes - size_pattern + 1 ;
    for ( j = 0 ; j < j_end ; j++ )
    {
        int distance = 0 ;
        int size ;

#ifdef APM_DEBUG
        if ( j % 10000 == 0 )
        {
            printf( "MPI rank %d : Processing byte %d (out of %d)\n", rank, j, part_bytes ) ;
            printf( "local matches of rank %d: ", rank ) ;
            for ( int i = 0 ; i < nb_patterns ; i++ )
            {
                printf( "%d, ", n_matches[i] ) ;
            }
            printf( "\n" ) ;
        }
#endif

        size = size_pattern ;
        // modifying the edge case for the last MPI process
        if ( part_bytes - j < size_pattern )
        {
            size = part_bytes - j ;
        }

        distance = levenshtein( pattern[i], &buf[j], size, column ) ;

        if ( distance <= approx_factor ) {
            n_matches[i]++ ;
        }
    }

    free( column ) ;
}
```



# MPI

## Main

```
/* Timer stop */
gettimeofday(&t2, NULL);

duration = (t2.tv_sec - t1.tv_sec) + ((t2.tv_usec - t1.tv_usec) / 1e6);

printf( "APM done in %lf s\n", duration ) ;

/*****
 * END MAIN LOOP
 *****/


for ( i = 0 ; i < nb_patterns ; i++ )
{
    printf( "Number of matches for pattern <S>: %d\n",
            pattern[i], n_matches[i] ) ;
}

return 0 ;
}

342+ /* Sum the matches of each process with a MPI reduction */
343+ MPI_Reduce(n_matches, global_matches, nb_patterns, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);
344+
345 /* Timer stop */
346 gettimeofday(&t2, NULL);
347
348 duration = (t2.tv_sec - t1.tv_sec) + ((t2.tv_usec - t1.tv_usec) / 1e6);
349+ #if APM_DEBUG
350+ printf("Rank %d finished :", rank);
351+ for ( i = 0 ; i < nb_patterns ; i++ ) {
352+     printf( "%d, ", n_matches[i] ) ;
353+ }
354+ printf("=====\n");
355+ #endif
356
357
358 /*****
359  * END MAIN LOOP
360  *****/
361
362+ if (rank == 0)
363+ {
364+     printf( "APM done in %lf s\n", duration ) ;
365+     for ( i = 0 ; i < nb_patterns ; i++ ) {
366+         printf( "Number of matches for pattern <S>: %d\n",
367+                 pattern[i], global_matches[i] ) ;
368+     }
369+ }
370+ MPI_Finalize();
371+ return 0 ;
372+ }
373
```

# MPI

## Trace



# MPI


## SpeedUp




# MPI

## SpeedUp

Average Speed up on MPI APM







INSTITUT  
POLYTECHNIQUE  
DE PARIS

# OpenMP+MPI



# OpenMP+MPI

## Main

```

106 int
107 main( int argc, char ** argv )
108 {
109     /* MPI initialisation
110     */
111     int nb_nodes;
112     int rank;
113     MPI_Status status;
114     MPI_Init(&argc, &argv);
115
116     MPI_Comm_size(MPI_COMM_WORLD, &nb_nodes);
117     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
118
119     #if APM_DEBUG
120     char hostname[256];
121     gethostname(hostname, sizeof(hostname));
122     printf("Process MPI rank %d of PID %d on %s ready for attach\n",rank, getpid(), hostname);
123     #endif
124
125     char ** pattern ;
126     char * filename ;
127     int approx_factor = 0 ;
128     int nb_patterns = 0 ;
129     int i, j ;
130     char * buf ;
131     struct timeval t1, t2;
132     double duration ;
133     int n_bytes ;
134
135     int * n_matches ;
136     int * global_matches;

```

```

106 int
107 main( int argc, char ** argv )
108 {
109     /* MPI initialisation
110     */
111     int nb_nodes;
112     int rank;
113     MPI_Status status;
114     int required = MPI_THREAD_FUNNELED;
115     int provided;
116     MPI_Init_thread(&argc, &argv, required, &provided);
117     MPI_Comm_size(MPI_COMM_WORLD, &nb_nodes);
118     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
119
120     #if APM_DEBUG
121     char hostname[256];
122     gethostname(hostname, sizeof(hostname));
123     printf("Process MPI rank %d of PID %d on %s ready for attach\n",rank, getpid(), hostname);
124     #endif
125
126     char ** pattern ;
127     char * filename ;
128     int approx_factor = 0 ;
129     int nb_patterns = 0 ;
130     int i, j ;
131     char * buf ;
132     struct timeval t1, t2;
133     double duration ;
134     int n_bytes ;
135     int tmp_matches ;
136     int * n_matches ;
137     int * global_matches;
138

```

# OpenMP+MPI

## Main

```
for ( i = 0 ; i < nb_patterns ; i++ )
{
    int size_pattern = strlen(pattern[i]) ;

    int * column ;

    n_matches[i] = 0 ;

    column = (int *)malloc( (size_pattern+1) * sizeof( int ) ) ;
    if ( column == NULL )
    {
        fprintf( stderr, "Error: unable to allocate memory for column (%ldB)\n",
            (size_pattern+1) * sizeof( int ) ) ;
        return 1 ;
    }

    int j_end = (rank == nb_nodes-1) ? part_bytes : part_bytes - size_pattern + 1 ;
    for ( j = 0 ; j < j_end ; j++ )
    {
        int distance = 0 ;
        int size ;
```

```
293 for ( i = 0 ; i < nb_patterns ; i++ )
294 {
295     int size_pattern = strlen(pattern[i]) ;
296
297     int * column ;
298
299     n_matches[i] = 0 ;
300
301+    tmp_matches = 0 ;
302+
303+    #pragma omp parallel
304+    {
305+        int j_end = (rank == nb_nodes-1) ? part_bytes : part_bytes - size_pattern + 1 ;
306+        #pragma omp for schedule(guided) reduction(+:tmp_matches)
307+        for ( j = 0 ; j < j_end ; j++ )
308+        {
309            column = (int *)malloc( (size_pattern+1) * sizeof( int ) ) ;
310+            if ( column == NULL ) {
311                fprintf( stderr, "Error: unable to allocate memory for column (%ldB)\n",
312                    (size_pattern+1) * sizeof( int ) ) ;
313+                exit(1);
314            }
315
316            int distance = 0 ;
317            int size ;
318
```

# OpenMP+MPI

## Main

```

size = size_pattern ;
// modifying the edge case for the last MPI process
if ( part_bytes - j < size_pattern )
{
    size = part_bytes - j ;
}

distance = levenshtein( pattern[i], buf[j], size, column ) ;

if ( distance <= approx_factor ) {
    n_matches[i]++ ;
}
}

free( column );
}

```

```

333 size = size_pattern ;
334 // modifying the edge case for the last MPI process
335 if ( part_bytes - j < size_pattern )
336 {
337     size = part_bytes - j ;
338 }
339
340 distance = levenshtein( pattern[i], &buf[j], size, column ) ;
341
342 if ( distance <= approx_factor ) {
343+     tmp_matches = tmp_matches + 1 ;
344+ }
345 }
346 }
347
348 free( column );
349+ n_matches[i] = tmp_matches;
350 }
351


```




# OpenMP+MPI

## Speed up


MPI and OMP APM with small pattern



MPI and OMP APM with medium pattern



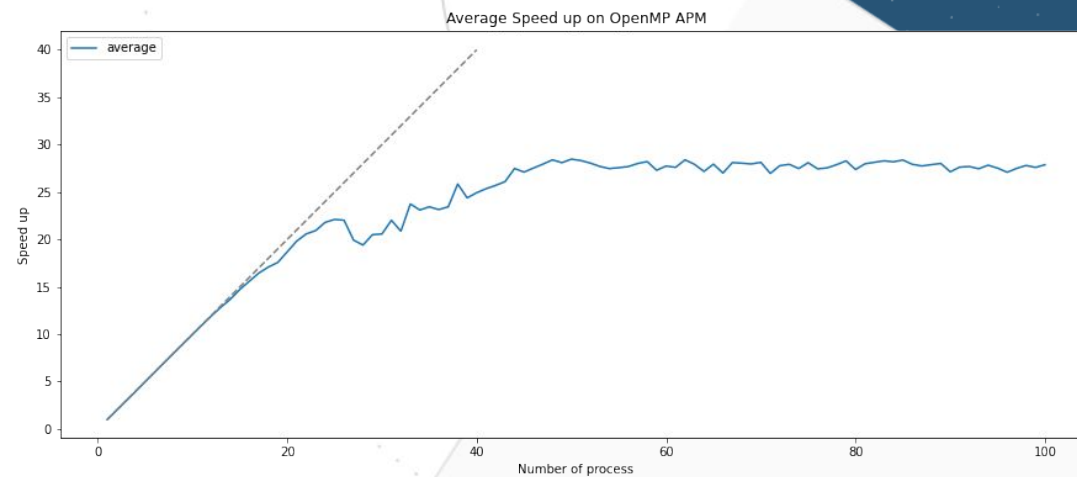

MPI and OMP APM with large pattern




# OpenMP+MPI

## Speed up

MPI and OMP APM with large pattern





INSTITUT  
POLYTECHNIQUE  
DE PARIS

# CUDA



# CUDA

## Idée générale

**ADN:**  
ACGT

**Principe:**

Calcul départ de A sur un GPU

Calcul départ de C sur un GPU

...

**Problème:**

Chaque GPU doit avoir en mémoire le pattern, la séquence d'ADN, ... (les paramètres en général) et un paramètre sur lequel stocker la distance à récupérer.

# CUDA

## Main

```

/* timer start */
gettimeofday(&t1, NULL);

for ( i = 0 ; i < nb_patterns ; i++ )
{
    int size_pattern = strlen(pattern[i]) ;

    int *column ;

    n_matches[i] = 0 ;

    column = (int *)malloc( (size_pattern+1) * sizeof( int ) ) ;
    if ( column == NULL )
    {
        fprintf( stderr, "Error: unable to allocate memory for column (%ldB)\n",
                (size_pattern+1) * sizeof( int ) ) ;
        return 1 ;
    }

    for ( j = 0 ; j < n_bytes ; j++ )
    {
        int distance = 0 ;
        int size ;

        #if APM_DEBUG
        if ( j % 100 == 0 )
        {
            printf( "Processing byte %d (out of %d)\n", j, n_bytes ) ;
        }
        #endif

        size = size_pattern ;
        if ( n_bytes - j < size_pattern )
        {
            size = n_bytes - j ;
        }

        distance = levenshtein( pattern[i], buf[j], size, column ) ;

        if ( distance <= approx_factor ) {
            n_matches[i]++ ;
        }
    }

    free( column ) ;
}

217 /* timer start */
218 gettimeofday(&t1, NULL);
219+ int blocksize = 1024;
220+ int nb_threads = min(blocksize,n_bytes) * ceil((n_bytes / (float)blocksize));
221+ dim3 dimBlock(min(blocksize,n_bytes));
222+ dim3 dimGrid(ceil((n_bytes / (float)blocksize)));
223
224 for ( i = 0 ; i < nb_patterns ; i++ )
225 {
226     int size_pattern = strlen(pattern[i]) ;
227
228+ int *nb_matches = (int *)malloc((n_bytes) * sizeof(int));
229+
230+ for (int j = 0; j < n_bytes ; j++ )
231+ {
232+     nb_matches[j] = 0;
233+ }
234
235     n_matches[i] = 0 ;| You, seconds ago - Uncommitted changes
236
237+ char *gpu_pattern;
238+ char *gpu_buf;
239+ int *gpu_matches;
240+ int *gpu_column;
241
242+ cudaMalloc(&gpu_pattern, (size_pattern) * sizeof(char));
243+ cudaMalloc(&gpu_buf, (n_bytes) * sizeof(char));
244+ cudaMalloc(&gpu_matches, (n_bytes) * sizeof(int));
245+ cudaMalloc(&gpu_column, nb_threads * (size_pattern + 1) * sizeof(int));
246
247+ cudaMemcpy(gpu_pattern, pattern[i], (size_pattern) * sizeof(char), cudaMemcpyHostToDevice);
248+ cudaMemcpy(gpu_buf, buf, (n_bytes) * sizeof(char), cudaMemcpyHostToDevice);
249+ cudaMemcpy(gpu_matches, nb_matches, (n_bytes) * sizeof(int), cudaMemcpyHostToDevice);
250
251+ cudaLevenshtein<<<dimGrid, dimBlock>>>(gpu_pattern, gpu_buf, size_pattern, n_bytes, approx_factor, gpu_column, gpu_matches);
252
253+ cudaMemcpy(nb_matches, gpu_matches, (n_bytes) * sizeof(int), cudaMemcpyDeviceToHost);
254
255+ for (int j = 0; j < n_bytes; j++)
256+ {
257+     n_matches[i] += nb_matches[j];
258+ }
259
260+ cudaFree(gpu_pattern);
261+ cudaFree(gpu_buf);
262+ cudaFree(gpu_matches);
263+ cudaFree(gpu_column);
264
265

```



# CUDA

## Levenshtein

```
int levenshtein(char *s1, char *s2, int len, int *column) {
    unsigned int x, y, lastdiag, olddiag;

    for (y = 1; y <= len; y++)
    {
        column[y] = y;
    }

    for (x = 1; x <= len; x++) {
        column[0] = x;

        lastdiag = x-1;
        for (y = 1; y <= len; y++) {
            olddiag = column[y];
            column[y] = MIN3(
                column[y] + 1,
                column[y-1] + 1,
                lastdiag + (s1[y-1] == s2[x-1] ? 0 : 1)
            );
            lastdiag = olddiag;
        }
    }

    return(column[len]);
}

__global__ void cuda_levenshtein(char *gpu_pattern, char *gpu_buf, int size_pattern, int n_bytes, int approx_factor, int *gpu_column, int *gpu_matches)
{
    unsigned int x, y, lastdiag, olddiag;

    int i = blockIdx.x * blockDim.x + threadIdx.x;

    gpu_column = &gpu_column[i * (size_pattern + 1)];
    gpu_buf = &gpu_buf[i];

    if (1 < n_bytes)
    {
        int distance = 0;
        int size;
        size = size_pattern;
        if (n_bytes - i < size_pattern)
        {
            size = n_bytes - i;
        }

        for (y = 1; y <= size; y++)
        {
            gpu_column[y] = y;
        }

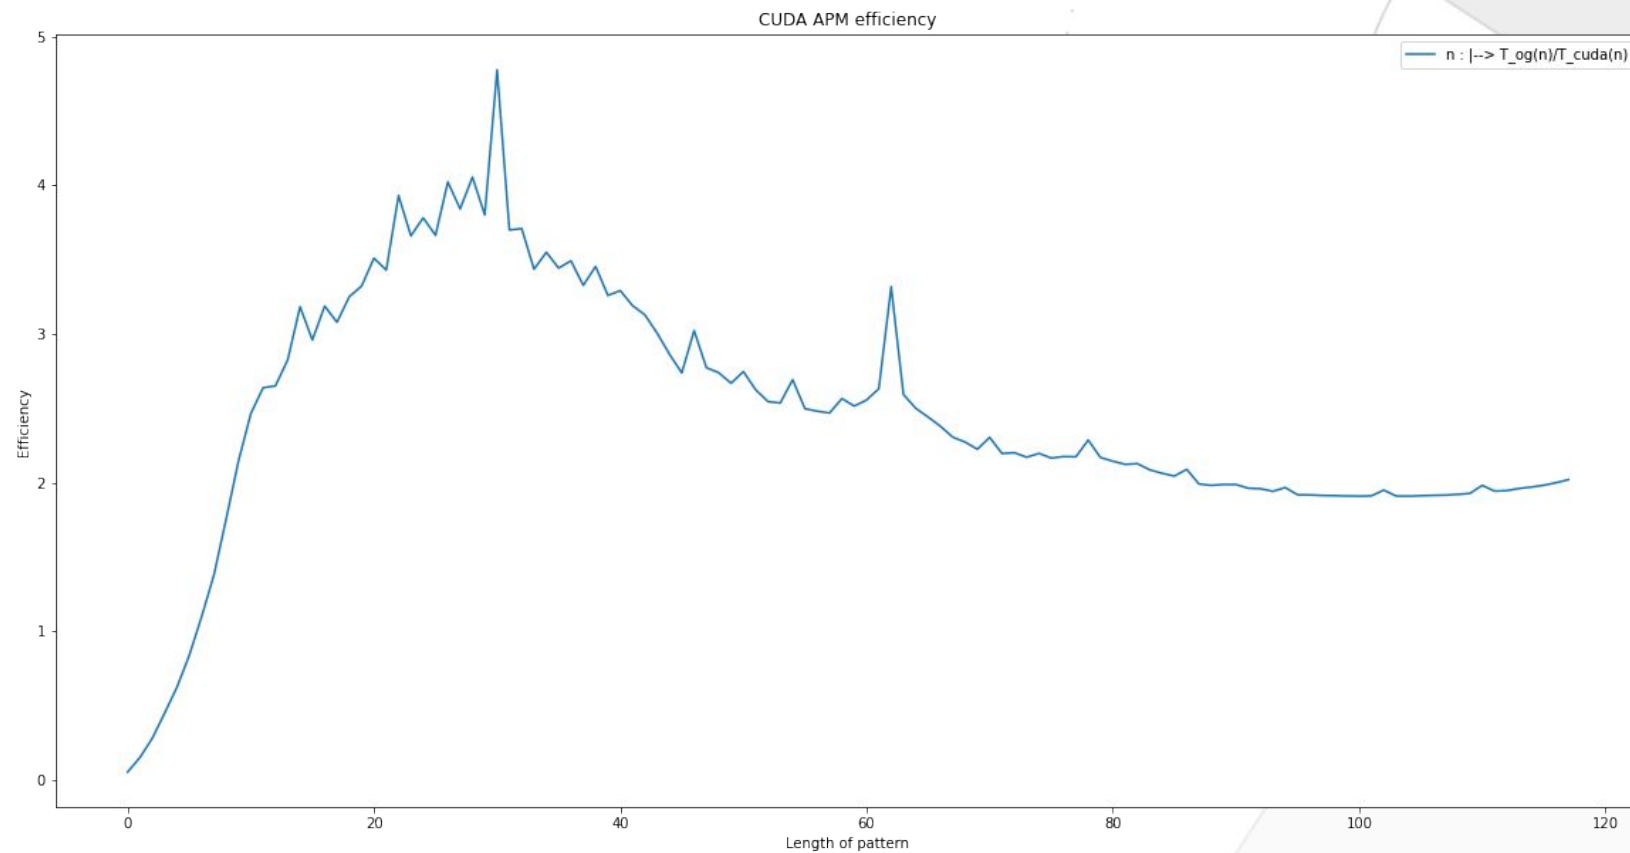
        for (x = 1; x <= size; x++)
        {
            gpu_column[0] = x;
            lastdiag = x-1;
            for (y = 1; y <= size; y++)
            {
                olddiag = gpu_column[y];
                gpu_column[y] = MIN3(
                    gpu_column[y] + 1,
                    gpu_column[y-1] + 1,
                    lastdiag + (gpu_pattern[y-1] == gpu_buf[x-1] ? 0 : 1)
                );
                lastdiag = olddiag;
            }
        }


        distance = gpu_column[size];

        if (distance <= approx_factor)
        {
            gpu_matches[i] = 1;
        }
        else
        {
            gpu_matches[i] = 0;
        }
    }
}
```

# CUDA

## Efficiency






INSTITUT  
POLYTECHNIQUE  
DE PARIS

# Comparaison




# Comparaison

SpeedUp OpenMp et MPI



# Comparaison

## SpeedUp OpenMp et MPI



MPI and OMP APM with large pattern

